

ch  PyTorch  PyTorch  PyTorch  PyTorch  PyTorch  PyTorch

ch  PyTorch  PyTorch  PyTorch  PyTorch  PyTorch  PyTorch

ch  PyTorch  PyTorch  PyTorch  PyTorch  PyTorch  PyTorch

ch  PyTorch  PyTorch  PyTorch  PyTorch  PyTorch  PyTorch

ch  PyTorch  PyTorch  PyTorch  PyTorch  PyTorch  PyTorch

ch  PyTorch  PyTorch  PyTorch  PyTorch  PyTorch  PyTorch

ch  PyTorch  PyTorch  PyTorch  PyTorch  PyTorch  PyTorch

ch  PyTorch  PyTorch  PyTorch  PyTorch  PyTorch  PyTorch

ch  PyTorch  PyTorch  PyTorch  PyTorch  PyTorch  PyTorch

ch  PyTorch  PyTorch  PyTorch  PyTorch  PyTorch  PyTorch

ch  PyTorch  PyTorch  PyTorch  PyTorch  PyTorch  PyTorch

ch  PyTorch  PyTorch  PyTorch  PyTorch  PyTorch  PyTorch

ch  PyTorch  PyTorch  PyTorch  PyTorch  PyTorch  PyTorch

ch  PyTorch  PyTorch  PyTorch  PyTorch  PyTorch  PyTorch

ch  PyTorch  PyTorch  PyTorch  PyTorch  PyTorch  PyTorch

# Usable while performant: the challenges building PyTorch

Soumith Chintala



# Problem Statement

- Deep Learning Workloads



# Problem Statement

- Deep Learning Workloads

```
for epoch in range(max_epochs):  
    for data, target in enumerate(training_data):  
        output = model(data)  
        loss = F.nll_loss(output, target)  
        loss.backward()  
        optimizer.step()
```

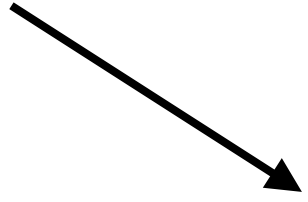


# Problem Statement

- Deep Learning Workloads

N samples, each of some shape D

```
for epoch in range(max_epochs):  
    for data, target in enumerate(training_data):  
        output = model(data)  
        loss = F.nll_loss(output, target)  
        loss.backward()  
        optimizer.step()
```



# Problem Statement

- Deep Learning Workloads mini-batch of  $M$  samples ( $M \ll N$ ), each of shape  $D$

```
for epoch in range(max_epochs):  
    for data, target in enumerate(training_data):  
        output = model(data)  
        loss = F.nll_loss(output, target)  
        loss.backward()  
        optimizer.step()
```



# Problem Statement

- Deep Learning Workloads

neural network with weights

```
for epoch in range(max_epochs) :  
    for data, target in enumerate(training_data) :  
        output = model(data)  
        loss = F.nll_loss(output, target)  
        loss.backward()  
        optimizer.step()
```



# Problem Statement

- Deep Learning Workloads      backpropagation:  
compute derivatives wrt loss, using chain rule

```
for epoch in range(max_epochs):  
    for data, target in enumerate(training_data):  
        output = model(data)  
        loss = F.nll_loss(output, target)  
        loss.backward()  
        optimizer.step()
```






# Problem Statement

- Deep Learning Workloads

update weights using the computed gradients

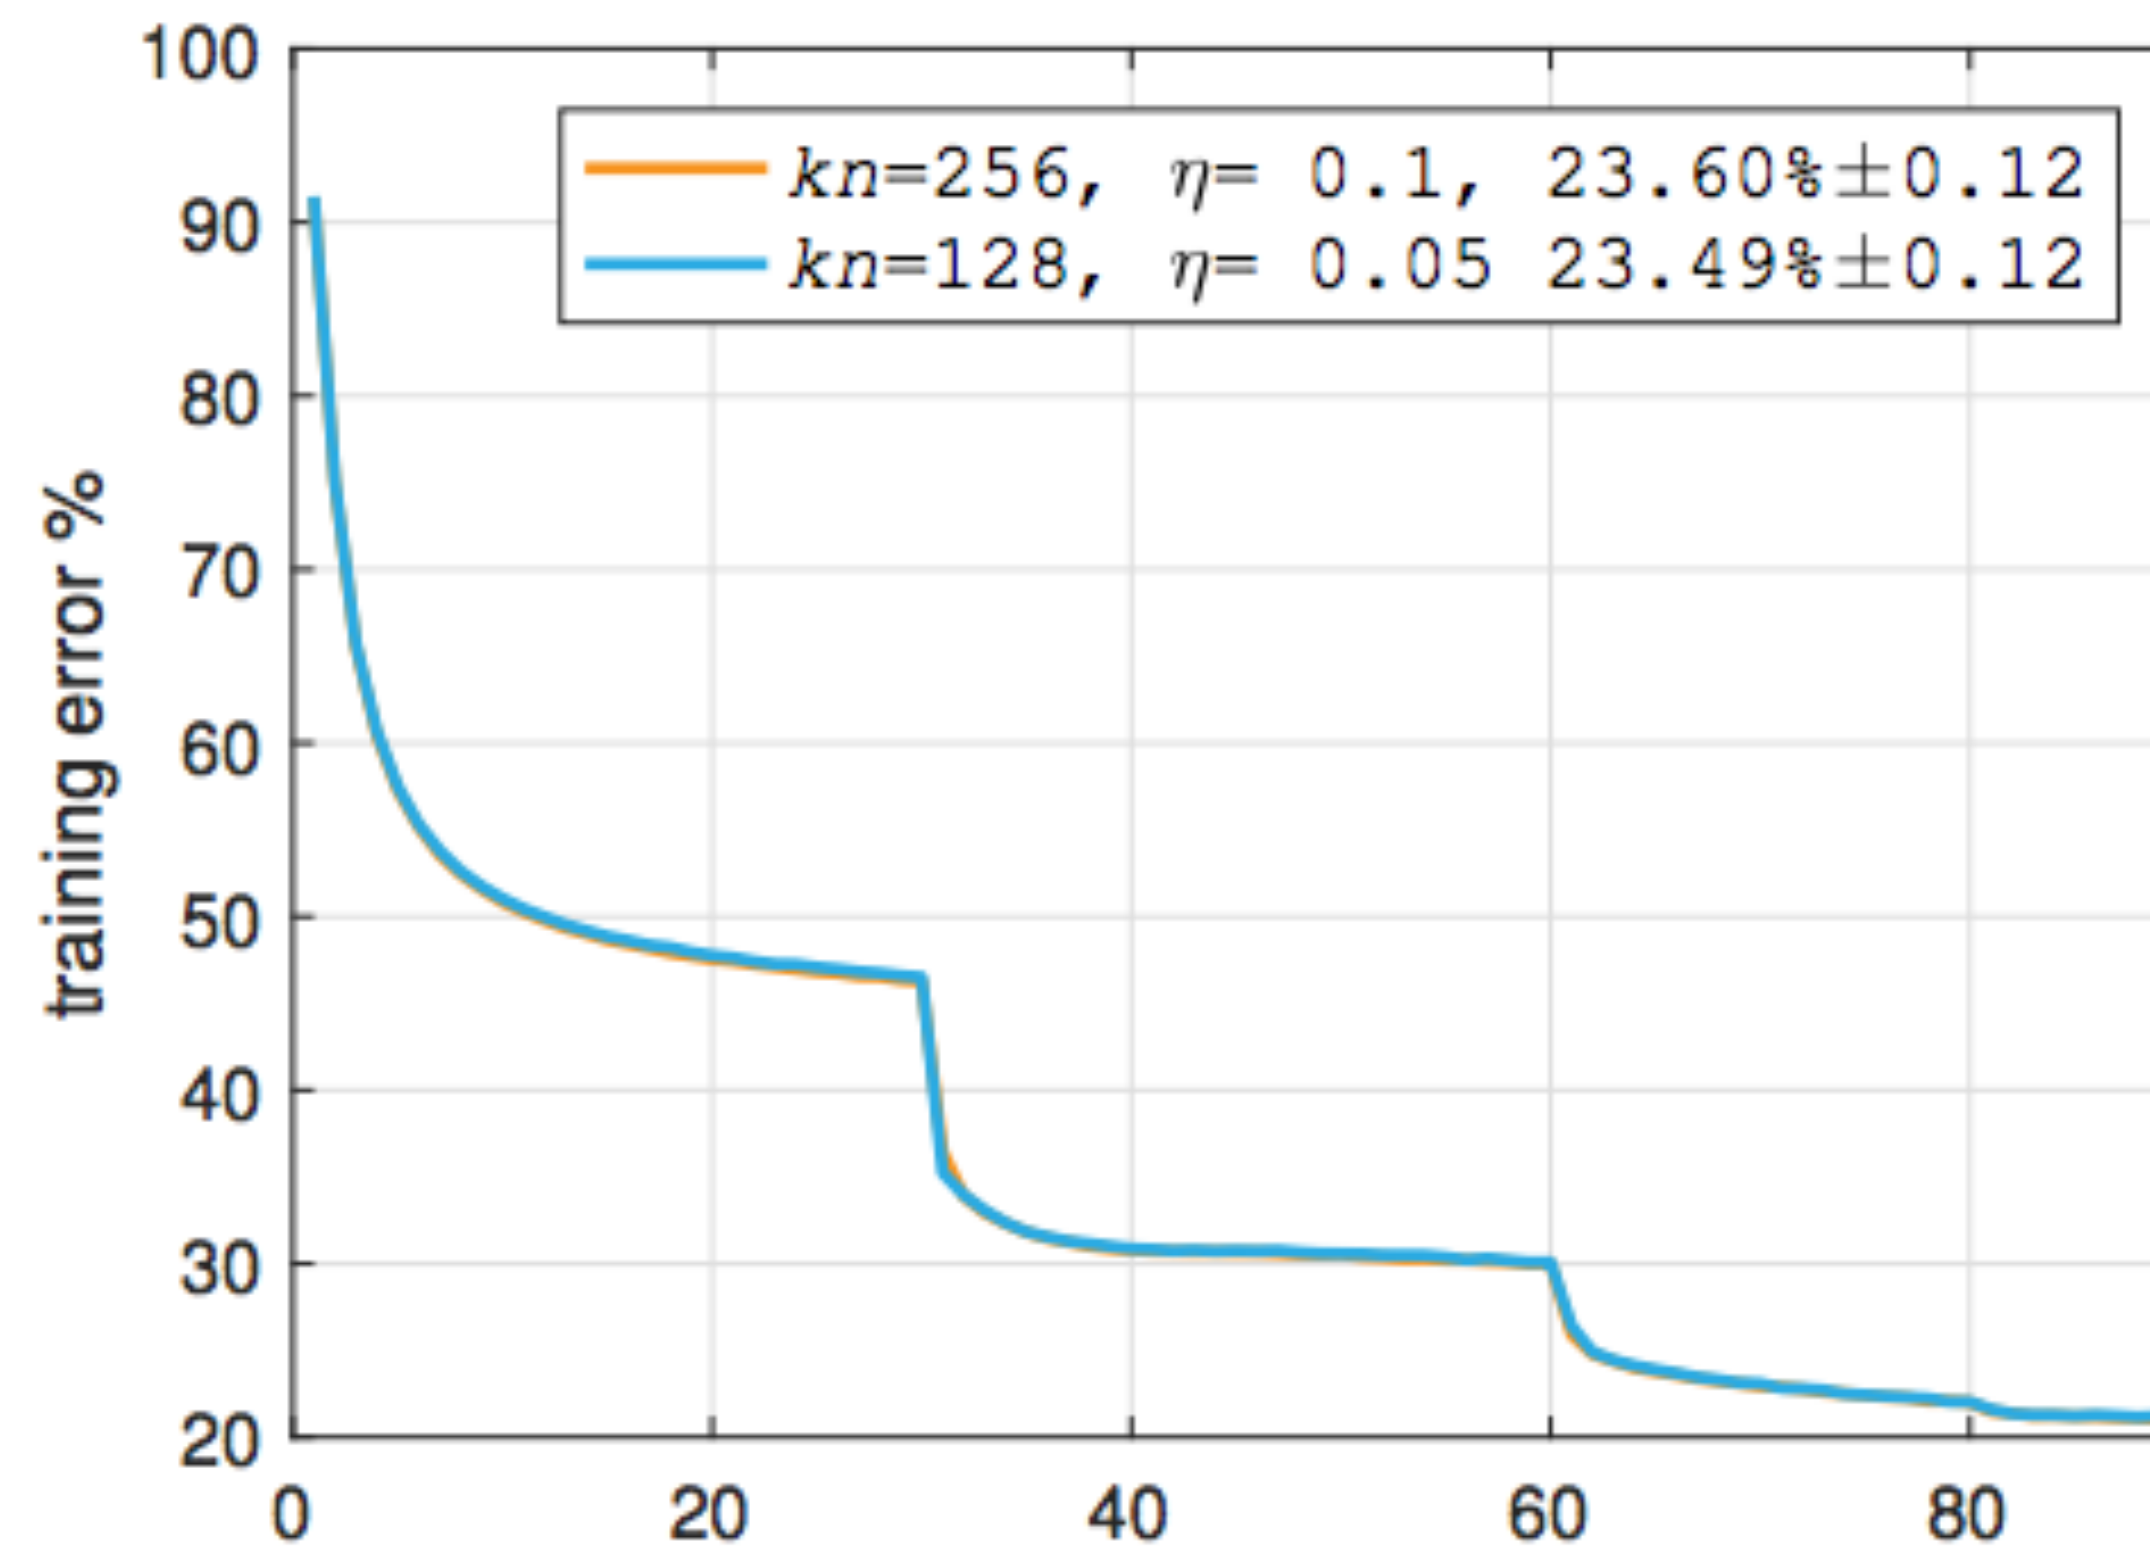
```
for epoch in range(max_epochs):  
    for data, target in enumerate(training_data):  
        output = model(data)  
        loss = F.nll_loss(output, target)  
        loss.backward()  
        optimizer.step()
```



# Problem Statement

- Deep Learning Workloads

```
for epoch  
for
```



```
.ning_data) :  
ret)
```



# Problem Statement

- Deep Learning Workloads

neural network with weights

```
for epoch in range(max_epochs) :  
    for data, target in enumerate(training_data) :  
        output = model(data)  
        loss = F.nll_loss(output, target)  
        loss.backward()  
        optimizer.step()
```



# Types of typical operators

## Convolution




Figure by Vincent Dumolin: [https://github.com/vdumoulin/conv\\_arithmetic](https://github.com/vdumoulin/conv_arithmetic)





# Types of typical operators

## Convolution






Figure by Vincent Dumolin: [https://github.com/vdumoulin/conv\\_arithmetic](https://github.com/vdumoulin/conv_arithmetic)



# Types of typical operators

## Convolution

```
for oc in output_channel:
    for ic in input_channel:
        for h in output_height:
            for w in output_width:
                for kh in kernel_height:
                    for kw in kernel_width:
                        output_pixel += input_pixel * kernel_value
```





# Types of typical operators






Figure by Vincent Dumolin: [https://github.com/vdumoulin/conv\\_arithmetic](https://github.com/vdumoulin/conv_arithmetic)



# Types of typical operators

## Matrix Multiply





# Types of typical operators

## Pointwise operations

```
for (i=0; i < data_length; i++) {  
    output[i] = input1[i] + input2[i]  
}
```



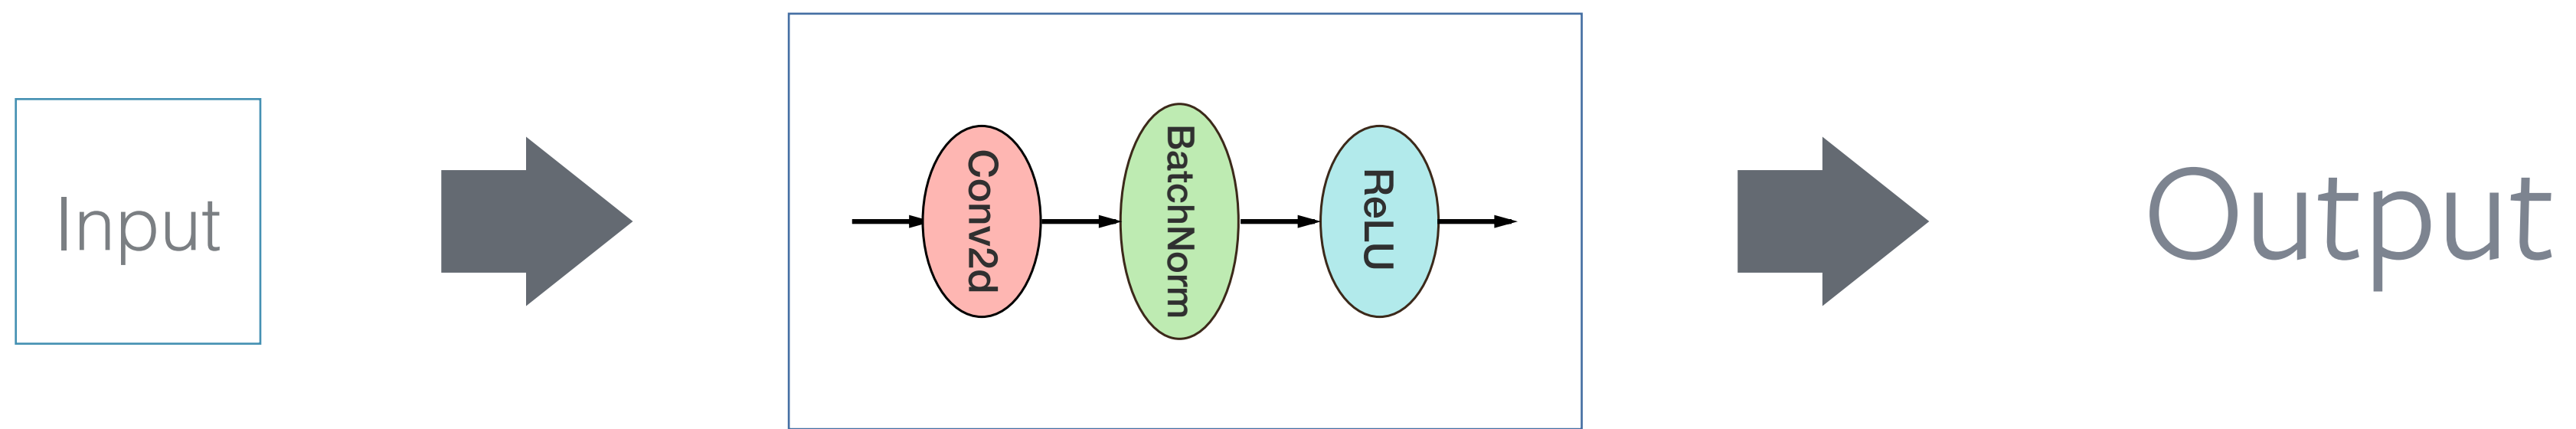
# Types of typical operators

## Reduction operations


```
double sum = 0.0;
for (i=0; i < data_length; i++) {
    sum += input[i];
}
```



# Chained Together

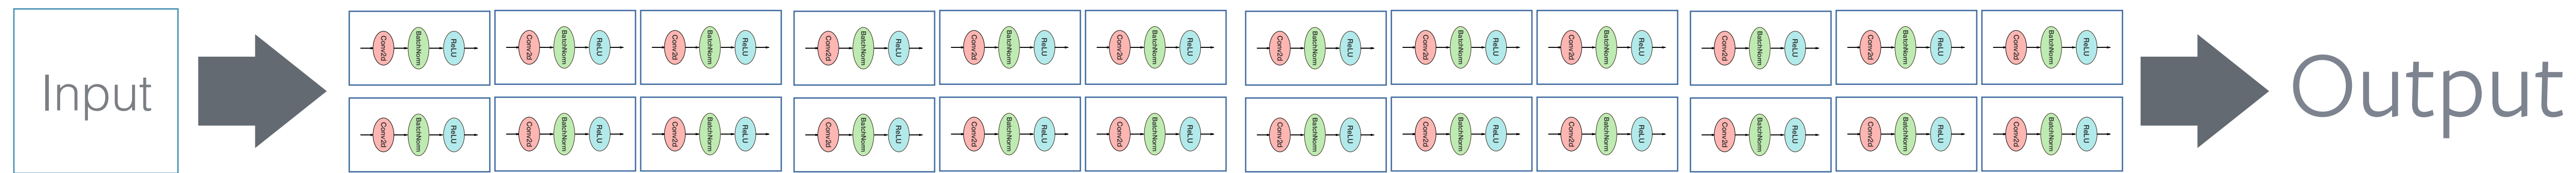


# Chained Together



# Chained Together


“deep”



# Chained Together

“deep”


recurrent



# Trained with Gradient Descent

“deep”

recurrent



# Problem Statement

- Deep Learning Workloads

an easy way to see recurrence

```
for epoch in range(max_epochs):  
    for data, target in enumerate(training_data):  
        output = model(data)  
        loss = F.nll_loss(output, target)  
        loss.backward()  
        optimizer.step()
```





# Problem Statement

- Deep Learning Workloads

an easy way to see recurrence

```
for epoch in range(max_epochs):  
    for data, target in enumerate(training_data):  
        output, hidden = [], zeros()  
        for t in data.size(0):  
            out, hidden = model(data[t], hidden)  
            output.append(out)  
        loss = F.nll_loss(output, target)  
        loss.backward()  
        optimizer.step()
```



# Problem Statement

- Deep Learning Workloads
  - Vision models
    - model is very deep, straight-line chain with no recurrence
    - lots of convolutions
    - typically run on GPUs



# Problem Statement

- Deep Learning Workloads

- Vision models
  - model is very deep, straight-line chain with no recurrence
  - lots of convolutions
  - typically run on GPUs
- NLP models
  - LSTM-RNN
    - model is 1 to 4 “layers” deep
    - two matmuls across space and time along with pointwise ops
    - typically run on CPUs if small, GPUs if large



# Deep Learning Frameworks

- Make this easy to program

```
for epoch in range(max_epochs):  
    for data, target in enumerate(training_data):  
        output = model(data)  
        loss = F.nll_loss(output, target)  
        loss.backward()  
        optimizer.step()
```



# Pre-PyTorch

Caffe

theano








meta programming meta programming

imperative



# Caffe

 <a href="#">deploy.prototxt</a>	[examples] switch examples + models to Input layers	3 years ago
 <a href="#">readme.md</a>	BVLC -> BAIR	2 years ago
 <a href="#">solver.prototxt</a>	Renaming CaffeNet model prototxts and unignoring models/*	4 years ago
 <a href="#">train_val.prototxt</a>	Upgrade existing nets using upgrade_net_proto_text tool	4 years ago

 <a href="#">readme.md</a>				
name	caffemodel	caffemodel_url	license	
BAIR/BVLC AlexNet Model	bvlc_alexnet.caffemodel	<a href="http://dl.caffe.berkeleyvision.org/bvlc_alexnet.caffemodel">http://dl.caffe.berkeleyvision.org/bvlc_alexnet.caffemodel</a>	unrestricted	91

```
layer {
  name: "conv1"
  type: "Convolution"
  bottom: "data"
  top: "conv1"
  param {
    lr_mult: 1
    decay_mult: 1
  }
  param {
    lr_mult: 2
    decay_mult: 0
  }
  convolution_param {
    num_output: 96
    kernel_size: 11
    stride: 4
    weight_filler {
      type: "gaussian"
      std: 0.01
    }
    bias_filler {
      type: "constant"
      value: 0
    }
  }
}
layer {
  name: "relu1"
  type: "ReLU"
  bottom: "conv1"
  top: "conv1"
}
layer {
  name: "norm1"
  type: "LRN"
  bottom: "conv1"
  top: "norm1"
  lrn_param {
    local_size: 5
    alpha: 0.0001
    beta: 0.75
  }
}
```



# Caffe

<a href="#">deploy.prototxt</a>	[examples] switch examples + models to Input layers	3 years ago
<a href="#">readme.md</a>	BVLC -> BAIR	2 years ago
<a href="#">solver.prototxt</a>	Renaming CaffeNet model prototxts and unignoring models/*	4 years ago
<a href="#">train_val.prototxt</a>		

readme.md				
name	caffemodel	caffemodel_url	license	
BAIR/BVLC AlexNet Model	bvlc_alexnet.caffemodel	<a href="http://dl.caffe.berkeleyvision.org/bvlc_alexnet.caffemodel">http://dl.caffe.berkeleyvision.org/bvlc_alexnet.caffemodel</a>	unrestricted	91

define protobuf, run via command-line utility

```
layer {
  name: "conv1"
  type: "Convolution"
  bottom: "data"
  top: "conv1"
  param {
    lr_mult: 1
    decay_mult: 1
  }
  param {
    lr_mult: 2
    decay_mult: 0
  }
  convolution_param {
    num_output: 96

    type: "gaussian"
    std: 0.01
  }
  bias_filler {
    type: "constant"
    value: 0
  }
}

layer {
  name: "relu1"
  type: "ReLU"
  bottom: "conv1"
  top: "conv1"
}

layer {
  name: "norm1"
  type: "LRN"
  bottom: "conv1"
  top: "norm1"
  lrn_param {
    local_size: 5
    alpha: 0.0001
    beta: 0.75
  }
}
```



# Caffe

<a href="#">deploy.prototxt</a>	[examples] switch examples + models to Input layers	3 years ago
<a href="#">readme.md</a>	BVLC -> BAIR	2 years ago
<a href="#">solver.prototxt</a>	Renaming CaffeNet model prototxts and unignoring models/*	4 years ago
<a href="#">train_val.prototxt</a>		

readme.md				
name	caffemodel	caffemodel url	license	
BAIR/BVLC AlexNet Model				

define protobuf, run via command-line utility

small, efficient library. Could do convents well.

```
layer {
  name: "conv1"
  type: "Convolution"
  bottom: "data"
  top: "conv1"
  param {
    lr_mult: 1
    decay_mult: 1
  }
  param {
    lr_mult: 2
    decay_mult: 0
  }
  convolution_param {
    num_output: 96

    type: "gaussian"
    std: 0.01
  }
  bias_filler {
    ..
  }
}

layer {
  name: "relu1"
  type: "ReLU"
  bottom: "conv1"
  top: "conv1"
}

layer {
  name: "norm1"
  type: "LRN"
  bottom: "conv1"
  top: "norm1"
  lrn_param {
    local_size: 5
    alpha: 0.0001
    beta: 0.75
  }
}
```





# Theano

```
# ##### BUILD NETWORK #####
# allocate symbolic variables for the data
# 'rand' is a random array used for random cropping/mirroring of data
x = T.ftensor4('x')
y = T.lvector('y')
rand = T.fvector('rand')

print '... building the model'
self.layers = []
params = []
weight_types = []

if flag_data_layer:
    data_layer = DataLayer(input=x, image_shape=(3, 256, 256,
                                                    batch_size),
                           cropsiz=227, rand=rand, mirror=True,
                           flag_rand=config['rand_crop'])

    layer1_input = data_layer.output
else:
    layer1_input = x

convpool_layer1 = ConvPoolLayer(input=layer1_input,
                                image_shape=(3, 227, 227, batch_size),
                                filter_shape=(3, 11, 11, 96),
                                convstride=4, padsize=0, group=1,
                                poolsize=3, poolstride=2,
                                bias_init=0.0, lrn=True,
                                lib_conv=lib_conv,
                                )
self.layers.append(convpool_layer1)
params += convpool_layer1.params
weight_types += convpool_layer1.weight_type

convpool_layer2 = ConvPoolLayer(input=convpool_layer1.output,
                                image_shape=(96, 27, 27, batch_size),
                                filter_shape=(96, 5, 5, 256),
                                convstride=1, padsize=2, group=2,
                                poolsize=3, poolstride=2,
                                bias_init=0.1, lrn=True,
                                lib_conv=lib_conv,
                                )
self.layers.append(convpool_layer2)
params += convpool_layer2.params
weight_types += convpool_layer2.weight_type
```

```
def compile_models(model, config, flag_top_5=False):
```

```
x = model.x
y = model.y
rand = model.rand
weight_types = model.weight_types
```

```
cost = model.cost
params = model.params
errors = model.errors
errors_top_5 = model.errors_top_5
batch_size = model.batch_size
```

```
mu = config['momentum']
eta = config['weight_decay']
```

```
# create a list of gradients for all model parameters
grads = T.grad(cost, params)
updates = []
```

```
learning_rate = theano.shared(np.float32(config['learning_rate']))
lr = T.scalar('lr') # symbolic learning rate
```

```
if config['use_data_layer']:
    raw_size = 256
else:
    raw_size = 227
```

```
shared_x = theano.shared(np.zeros((3, raw_size, raw_size,
                                     batch_size),
                                   dtype=theano.config.floatX),
                          borrow=True)
shared_y = theano.shared(np.zeros((batch_size,), dtype=int),
                          borrow=True)
```

```
rand_arr = theano.shared(np.zeros(3, dtype=theano.config.floatX),
                          borrow=True)
```

```
vels = [theano.shared(param_i.get_value() * 0.)
         for param_i in params]
```

```
if config['use_momentum']:
```

```
    assert len(weight_types) == len(params)
```

```
    for param_i, grad_i, vel_i, weight_type in \
```

```
# Define Theano Functions
```

```
train_model = theano.function([], cost, updates=updates,
                              givens=[(x, shared_x), (y, shared_y),
                                       (lr, learning_rate),
                                       (rand, rand_arr)])
```

```
validate_outputs = [cost, errors]
if flag_top_5:
    validate_outputs.append(errors_top_5)
```

```
validate_model = theano.function([], validate_outputs,
                                  givens=[(x, shared_x), (y, shared_y),
                                           (rand, rand_arr)])
```

```
train_error = theano.function(
    [], errors, givens=[(x, shared_x), (y, shared_y), (rand, rand_arr)])
```

```
return (train_model, validate_model, train_error,
        learning_rate, shared_x, shared_y, rand_arr, vels)
```



# Theano

```
# ##### BUILD NETWORK #####
# allocate symbolic variables for the data
# 'rand' is a random array used for random cropping/mirroring of data
x = T.ftensor4('x')
y = T.lvector('y')
rand = T.fvector('rand')

print '... building the model'
self.layers = []
params = []
weight_types = []

if flag_data_layer:
    data_layer = DataLayer(input=x, image_shape=(3, 256, 256,
                                                    batch_size),
                           cropsiz=227, rand=rand, mirror=True,

    layer1_input = c
else:
    layer1_input = x

convpool_layer1 = ConvPoolLayer(input=layer1_input,
                                image_shape=(3, 227, 227, batch_size),
                                filter_shape=(3, 11, 11, 96),
                                convstride=4, padsize=0, group=1,
                                poolsize=3, poolstride=2,
                                bias_init=0.0, lrn=True,
                                lib_conv=lib_conv,
                                )
self.layers.append(convpool_layer1)
params += convpool_layer1.params
weight_types += convpool_layer1.weight_type

convpool_layer2 = ConvPoolLayer(input=convpool_layer1.output,
                                image_shape=(96, 27, 27, batch_size),
                                filter_shape=(96, 5, 5, 256),
                                convstride=1, padsize=2, group=2,
                                poolsize=3, poolstride=2,
                                bias_init=0.1, lrn=True,
                                lib_conv=lib_conv,
                                )
self.layers.append(convpool_layer2)
params += convpool_layer2.params
weight_types += convpool_layer2.weight_type
```

```
def compile_models(model, config, flag_top_5=False):
```

```
x = model.x
y = model.y
rand = model.rand
weight_types = model.weight_types
```

```
cost = model.cost
params = model.params
errors = model.errors
errors_top_5 = model.errors_top_5
batch_size = model.batch_size
```

```
mu = config['momentum']
eta = config['weight_decay']
```

```
# create a list of gradients for all model parameters
grads = T.grad(cost, params)
updates = []
```

```
if config['use_data_layer']:
    raw_size = 256
else:
    raw_size = 227
```

```
shared_x = theano.shared(np.zeros((3, raw_size, raw_size,
                                     batch_size),
                                   dtype=theano.config.floatX),
                          borrow=True)
shared_y = theano.shared(np.zeros((batch_size,), dtype=int),
                          borrow=True)
```

```
rand_arr = theano.shared(np.zeros(3, dtype=theano.config.floatX),
                           borrow=True)
```

```
vels = [theano.shared(param_i.get_value() * 0.)
         for param_i in params]
```

```
if config['use_momentum']:
```

```
    assert len(weight_types) == len(params)
```

```
    for param_i, grad_i, vel_i, weight_type in \
```

```
# Define Theano Functions
```

```
train_model = theano.function([], cost, updates=updates,
                              givens=[(x, shared_x), (y, shared_y),
                                       (lr, learning_rate),
                                       (rand, rand_arr)])
```

```
validate_outputs = [cost, errors]
```

```
if flag_top_5:
    s_top_5)
    [], validate_outputs,
    givens=[(x, shared_x), (y, shared_y),
            (rand, rand_arr)])
```

```
train_error = theano.function(
    [], errors, givens=[(x, shared_x), (y, shared_y), (rand, rand_arr)])
```

```
return (train_model, validate_model, train_error,
        learning_rate, shared_x, shared_y, rand_arr, vels)
```

## meta-program Theano VM via Python API





# Theano

```
# ##### BUILD NETWORK #####
# allocate symbolic variables for the data
# 'rand' is a random array used for random cropping/mirroring of data
x = T.ftensor4('x')
y = T.lvector('y')
rand = T.fvector('rand')
```

```
print '... building the model'
self.layers = []
params = []
weight_types = []
```

```
if flag_data_layer:
    data_layer = DataLayer(input=x, image_shape=(3, 256, 256,
                                                    batch_size),
                           cropsiz=227, rand=rand, mirror=True,
```

```
    layer1_input = c
else:
    layer1_input = v
```

```
convpool_layer1 = ConvPoolLayer(input=layer1_input,
                                 filter_shape=(3, 11, 11, 96),
                                 convstride=4, padsize=0, group=1,
                                 poolsize=3, poolstride=2,
                                 bias_init=0.0, lrn=True,
                                 lib_conv=lib_conv,
                                 )
```

```
self.layers.append(convpool_layer1)
params += convpool_layer1.params
weight_types += convpool_layer1.weight_type
```

```
convpool_layer2 = ConvPoolLayer(input=convpool_layer1.output,
                                 image_shape=(96, 27, 27, batch_size),
                                 filter_shape=(96, 5, 5, 256),
                                 convstride=1, padsize=2, group=2,
                                 poolsize=3, poolstride=2,
                                 bias_init=0.1, lrn=True,
                                 lib_conv=lib_conv,
                                 )
```

```
self.layers.append(convpool_layer2)
params += convpool_layer2.params
weight_types += convpool_layer2.weight_type
```

```
def compile_models(model, config, flag_top_5=False):
```

```
    x = model.x
    y = model.y
    rand = model.rand
    weight_types = model.weight_types
```

```
    cost = model.cost
    params = model.params
    errors = model.errors
    errors_top_5 = model.errors_top_5
    batch_size = model.batch_size
```

```
    mu = config['momentum']
    eta = config['weight_decay']
```

```
    # create a list of gradients for all model parameters
    grads = T.grad(cost, params)
    updates = []
```

```
    if config['use_data_layer']:
```

```
    shared_x = theano.shared(np.zeros((3, raw_size, raw_size,
                                         batch_size),
                                         dtype=theano.config.floatX),
                              borrow=True)
```

```
    shared_y = theano.shared(np.zeros((batch_size,), dtype=int),
                              borrow=True)
```

```
    rand_arr = theano.shared(np.zeros(3, dtype=theano.config.floatX),
                              borrow=True)
```

```
    vels = [theano.shared(param_i.get_value() * 0.)
             for param_i in params]
```

```
    if config['use_momentum']:
```

```
        assert len(weight_types) == len(params)
```

```
        for param_i, grad_i, vel_i, weight_type in \
```

```
    # Define Theano Functions
```

```
    train_model = theano.function([], cost, updates=updates,
                                   givens=[(x, shared_x), (y, shared_y),
                                             (lr, learning_rate),
                                             (rand, rand_arr)])
```

```
    validate_outputs = [cost, errors]
```

```
    if flag_top_5:
```

```
        s_top_5)
```

```
    [], validate_outputs,
    givens=[(x, shared_x), (y, shared_y),
            (rand, rand_arr)])
```

```
    ), (y, shared_y), (rand, rand_arr)])
```

```
    return (train_model, validate_model, train_error,
            learning_rate, shared_x, shared_y, rand_arr, vels)
```

**meta-program Theano VM via Python API**  
**whole program optimizations, graph fusion**



# Theano

```
# ##### BUILD NETWORK #####
# allocate symbolic variables for the data
# 'rand' is a random array used for random cropping/mirroring of data
x = T.ftensor4('x')
y = T.lvector('y')
rand = T.fvector('rand')
```

```
print '... building the model'
self.layers = []
params = []
weight_types = []
```

```
if flag_data_layer:
    data_layer = DataLayer(input=x, image_shape=(3, 256, 256,
                                                    batch_size),
                           cropsize=227, rand=rand, mirror=True,
```

```
    layer1_input = c
else:
    layer1_input = v
```

```
convpool_layer1 = ConvPoolLayer(input=layer1_input,
                                 filter_shape=(3, 11, 11, 96),
```

```
    )
self.layers.append(convpool_layer1)
params += convpool_layer1.params
weight_types += convpool_layer1.weight_type
```

```
convpool_layer2 = ConvPoolLayer(input=convpool_layer1.output,
                                 image_shape=(96, 27, 27, batch_size),
                                 filter_shape=(96, 5, 5, 256),
                                 convstride=1, padsize=2, group=2,
                                 poolsize=3, poolstride=2,
                                 bias_init=0.1, lrn=True,
                                 lib_conv=lib_conv,
```

```
    )
self.layers.append(convpool_layer2)
params += convpool_layer2.params
weight_types += convpool_layer2.weight_type
```

```
def compile_models(model, config, flag_top_5=False):
```

```
    x = model.x
    y = model.y
    rand = model.rand
    weight_types = model.weight_types
```

```
    cost = model.cost
    params = model.params
    errors = model.errors
    errors_top_5 = model.errors_top_5
    batch_size = model.batch_size
```

```
    mu = config['momentum']
    eta = config['weight_decay']
```

```
    # create a list of gradients for all model parameters
    grads = T.grad(cost, params)
    updates = []
```

```
    if config['use_data_layer']:
```

```
        borrow=True,
        shared_y = theano.shared(np.zeros((batch_size,), dtype=int),
                                   borrow=True)
```

```
        rand_arr = theano.shared(np.zeros(3, dtype=theano.config.floatX),
                                   borrow=True)
```

```
        vels = [theano.shared(param_i.get_value() * 0.)
                  for param_i in params]
```

```
    if config['use_momentum']:
```

```
        assert len(weight_types) == len(params)
```

```
        for param_i, grad_i, vel_i, weight_type in \
```

```
    # Define Theano Functions
```

```
    train_model = theano.function([], cost, updates=updates,
                                   givens=[(x, shared_x), (y, shared_y),
                                             (lr, learning_rate),
                                             (rand, rand_arr)])
```

```
    validate_outputs = [cost, errors]
```

```
    if flag_top_5:
```

```
        s_top_5)
```

```
    [], validate_outputs,
    givens=[(x, shared_x), (y, shared_y),
            (rand, rand_arr)])
```

```
    ), (y, shared_y), (rand, rand_arr)])
```

```
    vels)
```

meta-program Theano VM via Python API

whole program optimizations, graph fusion

graphs took minutes to hours to compile and start





# Torch-7

```
function alexnet(lib)
    local SpatialConvolution = lib[1]
    local SpatialMaxPooling = lib[2]
    local ReLU = lib[3]
    local SpatialZeroPadding = nn.SpatialZeroPadding
    local padding = true
    local stride1only = false

    -- from https://code.google.com/p/cuda-convnet2/source/browse/layers/layers-imagenet-1gpu.cfg
    -- this is AlexNet that was presented in the One Weird Trick paper. http://arxiv.org/abs/1404.5997
    local features = nn.Sequential()
    features:add(SpatialConvolution(3,64,11,11,4,4,2,2))      -- 224 -> 55
    features:add(ReLU(true))
    features:add(SpatialMaxPooling(3,3,2,2))                -- 55 -> 27
    features:add(SpatialConvolution(64,192,5,5,1,1,2,2))    -- 27 -> 27
    features:add(ReLU(true))
    features:add(SpatialMaxPooling(3,3,2,2))                -- 27 -> 13
    features:add(SpatialConvolution(192,384,3,3,1,1,1,1))    -- 13 -> 13
    features:add(ReLU(true))
    features:add(SpatialConvolution(384,256,3,3,1,1,1,1))    -- 13 -> 13
    features:add(ReLU(true))
    features:add(SpatialConvolution(256,256,3,3,1,1,1,1))    -- 13 -> 13
    features:add(ReLU(true))
    features:add(SpatialMaxPooling(3,3,2,2))                -- 13 -> 6

    local classifier = nn.Sequential()
    classifier:add(nn.View(256*6*6))
    -- classifier:add(nn.Dropout(0.5))
    classifier:add(nn.Linear(256*6*6, 4096))
    classifier:add(nn.Threshold(0, 1e-6))
    -- classifier:add(nn.Dropout(0.5))
    classifier:add(nn.Linear(4096, 4096))
    classifier:add(nn.Threshold(0, 1e-6))
    classifier:add(nn.Linear(4096, 1000))
    -- classifier:add(nn.LogSoftMax())

    features:get(1).gradInput = nil

    local model = nn.Sequential()
    model:add(features):add(classifier)

    return model, 'AlexNet', {128, 3, 224, 224}
end

return alexnet
```

-- 4. trainBatch - Used by train() to train a single batch after the data is loaded.

```
function trainBatch(inputsCPU, labelsCPU)
    cutorch.synchronize()
    collectgarbage()
    local dataLoadingTime = dataTimer:time().real
    timer:reset()

    -- transfer over to GPU
    inputs:resize(inputsCPU:size()):copy(inputsCPU)
    labels:resize(labelsCPU:size()):copy(labelsCPU)

    local err, outputs
    feval = function(x)
        model:zeroGradParameters()
        outputs = model:forward(inputs)
        err = criterion:forward(outputs, labels)
        local gradOutputs = criterion:backward(outputs, labels)
        model:backward(inputs, gradOutputs)
        return err, gradParameters
    end
    optim.sgd(feval, parameters, optimState)

    cutorch.synchronize()
    batchNumber = batchNumber + 1
    loss_epoch = loss_epoch + err
    -- top-1 error
    local top1 = 0
    do
        local _,prediction_sorted = outputs:float():sort(2, true) -- descending
        for i=1,opt.batchSize do
            if prediction_sorted[i][1] == labelsCPU[i] then
                top1_epoch = top1_epoch + 1;
                top1 = top1 + 1
            end
        end
        top1 = top1 * 100 / opt.batchSize;
    end
    -- Calculate top-1 error, and print information
    print(('Epoch: [%d][%d/%d]\tTime %.3f Err %.4f Top1-%=: %.2f LR %.0e DataLoadingTime %.3f'):format(
        epoch, batchNumber, opt.epochSize, timer:time().real, err, top1,
        optimState.learningRate, dataLoadingTime))

    dataTimer:reset()
end
```

# Torch-7

```
function alexnet(lib)
    local SpatialConvolution = lib[1]
    local SpatialMaxPooling = lib[2]
    local ReLU = lib[3]
    local SpatialZeroPadding = nn.SpatialZeroPadding
    local padding = true
    local stride1only = false

    -- from https://code.google.com/p/cuda-convnet2/source/browse/layers/layers-imagenet-1gpu.cfg
    -- this is AlexNet that was presented in the One Weird Trick paper. http://arxiv.org/abs/1404.5997
    local features = nn.Sequential()
    features:add(SpatialConvolution(3,64,11,11,4,4,2,2)) -- 224 -> 55
    features:add(ReLU(true))
    features:add(SpatialMaxPooling(3,3,2,2)) -- 55 -> 27
    features:add(SpatialConvolution(64,192,5,5,1,1,2,2)) -- 27 -> 27
    features:add(ReLU(true))
    features:add(SpatialMaxPooling(3,3,2,2)) -- 27 -> 13
    features:add(SpatialConvolution(128,128,5,5,1,1,2,2))
    features:add(ReLU(true))
    features:add(SpatialConvolution(128,128,5,5,1,1,2,2))
    features:add(ReLU(true))
    features:add(SpatialConvolution(128,128,5,5,1,1,2,2))
    features:add(ReLU(true))
    features:add(SpatialMaxPooling(3,3,2,2))

    local classifier = nn.Sequential()
    classifier:add(nn.View(256*6*6))
    -- classifier:add(nn.Dropout(0.5))
    classifier:add(nn.Linear(256*6*6, 4096))
    classifier:add(nn.Threshold(0, 1e-6))
    -- classifier:add(nn.Dropout(0.5))
    classifier:add(nn.Linear(4096, 4096))
    classifier:add(nn.Threshold(0, 1e-6))
    classifier:add(nn.Linear(4096, 1000))
    -- classifier:add(nn.LogSoftMax())

    features:get(1).gradInput = nil

    local model = nn.Sequential()
    model:add(features):add(classifier)

    return model, 'AlexNet', {128, 3, 224, 224}
end

return alexnet
```

## imperative programming in Lua

```
-- 4. trainBatch - Used by train() to train a single batch after the data is loaded.
function trainBatch(inputsCPU, labelsCPU)
    cutorch.synchronize()
    collectgarbage()
    local dataLoadingTime = dataTimer:time().real
    timer:reset()

    -- transfer over to GPU
    inputs:resize(inputsCPU:size()):copy(inputsCPU)
    labels:resize(labelsCPU:size()):copy(labelsCPU)

    local err, outputs
    feval = function(x)
        model:zeroGradParameters()
        outputs = model:forward(inputs)
        err = criterion:forward(outputs, labels)
        return err, outputs
    end

    loss_epoch = loss_epoch + err
    -- top-1 error
    local top1 = 0
    do
        local _,prediction_sorted = outputs:float():sort(2, true) -- descending
        for i=1,opt.batchSize do
            if prediction_sorted[i][1] == labelsCPU[i] then
                top1_epoch = top1_epoch + 1;
                top1 = top1 + 1
            end
        end
        top1 = top1 * 100 / opt.batchSize;
    end
    -- Calculate top-1 error, and print information
    print(('Epoch: [%d][%d/%d]\tTime %.3f Err %.4f Top1-%=: %.2f LR %.0e DataLoadingTime %.3f'):format(
        epoch, batchNumber, opt.epochSize, timer:time().real, err, top1,
        optimState.learningRate, dataLoadingTime))

    dataTimer:reset()
end
```



# Torch-7

```
function alexnet(lib)
    local SpatialConvolution = lib[1]
    local SpatialMaxPooling = lib[2]
    local ReLU = lib[3]
    local SpatialZeroPadding = nn.SpatialZeroPadding
    local padding = true
    local stride1only = false

    -- from https://code.google.com/p/cuda-convnet2/source/browse/layers/layers-imagenet-1gpu.cfg
    -- this is AlexNet that was presented in the One Weird Trick paper. http://arxiv.org/abs/1404.5997
    local features = nn.Sequential()
    features:add(SpatialConvolution(3,64,11,11,4,4,2,2)) -- 224 -> 55
    features:add(ReLU(true))
    features:add(SpatialMaxPooling(3,3,2,2)) -- 55 -> 27
    features:add(SpatialConvolution(64,192,5,5,1,1,2,2)) -- 27 -> 27
```

**imperative programming in Lua**  
**tied closely to underlying C89 implementations**

```
classifier:add(nn.ReLU(true))
-- classifier:add(nn.Dropout(0.5))
classifier:add(nn.Linear(256*6*6, 4096))
classifier:add(nn.Threshold(0, 1e-6))
-- classifier:add(nn.Dropout(0.5))
classifier:add(nn.Linear(4096, 4096))
classifier:add(nn.Threshold(0, 1e-6))
classifier:add(nn.Linear(4096, 1000))
-- classifier:add(nn.LogSoftMax())

features:get(1).gradInput = nil

local model = nn.Sequential()
model:add(features):add(classifier)

return model, 'AlexNet', {128, 3, 224, 224}
end
```

```
return alexnet
```

-- 4. trainBatch - Used by train() to train a single batch after the data is loaded.

```
function trainBatch(inputsCPU, labelsCPU)
    cutorch.synchronize()
    collectgarbage()
    local dataLoadingTime = dataTimer:time().real
    timer:reset()
```

```
-- transfer over to GPU
inputs:resize(inputsCPU:size()):copy(inputsCPU)
labels:resize(labelsCPU:size()):copy(labelsCPU)
```

```
local err, outputs
feval = function(x)
    model:zeroGradParameters()
```

```
local top1 = 0
do
    local _,prediction_sorted = outputs:float():sort(2, true) -- descending
    for i=1,opt.batchSize do
        if prediction_sorted[i][1] == labelsCPU[i] then
            top1_epoch = top1_epoch + 1;
            top1 = top1 + 1
        end
    end
    top1 = top1 * 100 / opt.batchSize;
end
-- Calculate top-1 error, and print information
print(('Epoch: [%d] [%d/%d]\tTime %.3f Err %.4f Top1-%=: %.2f LR %.0e DataLoadingTime %.3f'):format(
    epoch, batchNumber, opt.epochSize, timer:time().real, err, top1,
    optimState.learningRate, dataLoadingTime))
```

```
dataTimer:reset()
```

```
end
```

# Torch-7

```
function alexnet(lib)
    local SpatialConvolution = lib[1]
    local SpatialMaxPooling = lib[2]
    local ReLU = lib[3]
    local SpatialZeroPadding = nn.SpatialZeroPadding
    local padding = true
    local stride1only = false

    -- from https://code.google.com/p/cuda-convnet2/source/browse/layers/layers-imagenet-1gpu.cfg
    -- this is AlexNet that was presented in the One Weird Trick paper. http://arxiv.org/abs/1404.5997
    local features = nn.Sequential()
    features:add(SpatialConvolution(3,64,11,11,4,4,2,2))    -- 224 -> 55
    features:add(ReLU(true))
    features:add(SpatialMaxPooling(3,3,2,2))                -- 55 -> 27
    features:add(SpatialConvolution(64,192,5,5,1,1,2,2))    -- 27 -> 27
```

**imperative programming in Lua**  
**tied closely to underlying C89 implementations**  
**Lua lacked good tooling and ecosystem**

```
-- classifier: add(nn.ReLU(true))
-- classifier:add(nn.Dropout(0.5))
classifier:add(nn.Linear(256*6*6, 4096))
classifier:add(nn.Threshold(0, 1e-6))
-- classifier:add(nn.Dropout(0.5))
classifier:add(nn.Linear(4096, 4096))
classifier:add(nn.Threshold(0, 1e-6))
classifier:add(nn.Linear(4096, 1000))
-- classifier:add(nn.LogSoftMax())

features:get(1).gradInput = nil

local model = nn.Sequential()
model:add(features):add(classifier)

return model, 'AlexNet', {128, 3, 224, 224}
end
```

```
return alexnet
```

-- 4. trainBatch - Used by train() to train a single batch after the data is loaded.

```
function trainBatch(inputsCPU, labelsCPU)
    cutorch.synchronize()
    collectgarbage()
    local dataLoadingTime = dataTimer:time().real
    timer:reset()
```

```
-- transfer over to GPU
inputs:resize(inputsCPU:size()):copy(inputsCPU)
labels:resize(labelsCPU:size()):copy(labelsCPU)
```

```
local err, outputs
feval = function(x)
    model:zeroGradParameters()
```


```
    local top1 = 0
    do
        local _,prediction_sorted = outputs:float():sort(2, true) -- descending
        for i=1,opt.batchSize do
            if prediction_sorted[i][1] == labelsCPU[i] then
                top1_epoch = top1_epoch + 1;
                top1 = top1 + 1
            end
        end
        top1 = top1 * 100 / opt.batchSize;
    end
    -- Calculate top-1 error, and print information
    print(('Epoch: [%d][%d/%d]\tTime %.3f Err %.4f Top1-%s: %.2f LR %.0e DataLoadingTime %.3f'):format(
        epoch, batchNumber, opt.epochSize, timer:time().real, err, top1,
        optimState.learningRate, dataLoadingTime))
```

```
    dataTimer:reset()
```

```
end
```



# What is PyTorch?



Deep Learning

Numpy-alternative

Reinforcement Learning



# ndarray library

- `np.ndarray`  $\leftrightarrow$  `torch.Tensor`
- 200+ operations, similar to numpy
- very fast acceleration on NVIDIA GPUs



```
# -*- coding: utf-8 -*-
```

```
import numpy as np
```

```
# N is batch size; D_in is input dimension;
```

```
# H is hidden dimension; D_out is output dimension.
```

```
N, D_in, H, D_out = 64, 1000, 100, 10
```

```
# Create random input and output data
```

```
x = np.random.randn(N, D_in)
```

```
y = np.random.randn(N, D_out)
```

```
# Randomly initialize weights
```

```
w1 = np.random.randn(D_in, H)
```

```
w2 = np.random.randn(H, D_out)
```

```
learning_rate = 1e-6
```

```
for t in range(500):
```

```
    # Forward pass: compute predicted y
```

```
    h = x.dot(w1)
```

```
    h_relu = np.maximum(h, 0)
```

```
    y_pred = h_relu.dot(w2)
```

```
    # Compute and print loss
```

```
    loss = np.square(y_pred - y).sum()
```

```
    print(t, loss)
```

```
    # Backprop to compute gradients of w1 and w2 with respect to loss
```

```
    grad_y_pred = 2.0 * (y_pred - y)
```

```
    grad_w2 = h_relu.T.dot(grad_y_pred)
```

```
    grad_h_relu = grad_y_pred.dot(w2.T)
```

```
    grad_h = grad_h_relu.copy()
```

```
    grad_h[h < 0] = 0
```

```
    grad_w1 = x.T.dot(grad_h)
```

```
    # Update weights
```

```
    w1 -= learning_rate * grad_w1
```

```
    w2 -= learning_rate * grad_w2
```

# Numpy

```
import torch
```

```
dtype = torch.FloatTensor
```

```
# dtype = torch.cuda.FloatTensor # Uncomment this to run on GPU
```

```
# N is batch size; D_in is input dimension;
```

```
# H is hidden dimension; D_out is output dimension.
```

```
N, D_in, H, D_out = 64, 1000, 100, 10
```

```
# Create random input and output data
```

```
x = torch.randn(N, D_in).type(dtype)
```

```
y = torch.randn(N, D_out).type(dtype)
```

```
# Randomly initialize weights
```

```
w1 = torch.randn(D_in, H).type(dtype)
```

```
w2 = torch.randn(H, D_out).type(dtype)
```

```
learning_rate = 1e-6
```

```
for t in range(500):
```

```
    # Forward pass: compute predicted y
```

```
    h = x.mm(w1)
```

```
    h_relu = h.clamp(min=0)
```

```
    y_pred = h_relu.mm(w2)
```

```
    # Compute and print loss
```

```
    loss = (y_pred - y).pow(2).sum()
```

```
    print(t, loss)
```

```
    # Backprop to compute gradients of w1 and w2 with respect to loss
```

```
    grad_y_pred = 2.0 * (y_pred - y)
```

```
    grad_w2 = h_relu.t().mm(grad_y_pred)
```

```
    grad_h_relu = grad_y_pred.mm(w2.t())
```

```
    grad_h = grad_h_relu.clone()
```

```
    grad_h[h < 0] = 0
```

```
    grad_w1 = x.t().mm(grad_h)
```

```
    # Update weights using gradient descent
```

```
    w1 -= learning_rate * grad_w1
```

```
    w2 -= learning_rate * grad_w2
```

# PyTorch

# ndarray / Tensor library

Tensors are similar to numpy's ndarrays, with the addition being that Tensors can also be used on a GPU to accelerate computing.

```
from __future__ import print_function
import torch
```

Construct a 5x3 matrix, uninitialized:

```
x = torch.Tensor(5, 3)
print(x)
```

Out:

```
1.000000e-25 *
 0.4136  0.0000  0.0000
 0.0000  1.6519  0.0000
 1.6518  0.0000  1.6519
 0.0000  1.6518  0.0000
 1.6520  0.0000  1.6519
[torch.FloatTensor of size 5x3]
```



# ndarray / Tensor library

Construct a randomly initialized matrix

```
x = torch.rand(5, 3)
print(x)
```

Out:

```
0.2598  0.7231  0.8534
0.3928  0.1244  0.5110
0.5476  0.2700  0.5856
0.7288  0.9455  0.8749
0.6663  0.8230  0.2713
[torch.FloatTensor of size 5x3]
```

Get its size

```
print(x.size())
```

Out:

```
torch.Size([5, 3])
```





# ndarray / Tensor library

You can use standard numpy-like indexing with all bells and whistles!

```
print(x[:, 1])
```

Out:

```
0.7231  
0.1244  
0.2700  
0.9455  
0.8230  
[torch.FloatTensor of size 5]
```



# ndarray / Tensor library

```
y = torch.rand(5, 3)  
print(x + y)
```

Out:

```
0.7931  1.1872  1.6143  
1.1946  0.4669  0.9639  
0.7576  0.8136  1.1897  
0.7431  1.8579  1.3400  
0.8188  1.1041  0.8914  
[torch.FloatTensor of size 5x3]
```



# NumPy bridge

## Converting torch Tensor to numpy Array

```
a = torch.ones(5)  
print(a)
```

Out:

```
1  
1  
1  
1  
1  
[torch.FloatTensor of size 5]
```

```
b = a.numpy()  
print(b)
```

Out:

```
[ 1.  1.  1.  1.  1.]
```





# NumPy bridge

## Converting torch Tensor to numpy Array

```
a = torch.ones(5)  
print(a)
```

Out:

```
1  
1  
1  
1  
1  
[torch.FloatTensor of size 5]
```

**Zero memory-copy  
very efficient**

```
b = a.numpy()  
print(b)
```

Out:

```
[ 1.  1.  1.  1.  1.]
```



# NumPy bridge

See how the numpy array changed in value.

```
a.add_(1)  
print(a)  
print(b)
```

Out:

```
2  
2  
2  
2  
2  
[torch.FloatTensor of size 5]  
  
[ 2.  2.  2.  2.  2.]
```



# NumPy bridge

## Converting numpy Array to torch Tensor

See how changing the np array changed the torch Tensor automatically

```
import numpy as np
a = np.ones(5)
b = torch.from_numpy(a)
np.add(a, 1, out=a)
print(a)
print(b)
```

Out:

```
[ 2.  2.  2.  2.  2.]

2
2
2
2
2
[torch.DoubleTensor of size 5]
```

All the Tensors on the CPU except a CharTensor support converting to NumPy and back.



# Seamless GPU Tensors

## CUDA Tensors 🔗

Tensors can be moved onto GPU using the `.cuda` function.

```
# let us run this cell only if CUDA is available
if torch.cuda.is_available():
    x = x.cuda()
    y = y.cuda()
    x + y
```




# automatic differentiation engine

for deep learning and reinforcement learning



# PyTorch Autograd


```
W_h = torch.randn(20, 20, requires_grad=True)
W_x = torch.randn(20, 10, requires_grad=True)
x = torch.randn(1, 10)
prev_h = torch.randn(1, 20)
```



# PyTorch Autograd

```
W_h = torch.randn(20, 20, requires_grad=True)
W_x = torch.randn(20, 10, requires_grad=True)
x = torch.randn(1, 10)
prev_h = torch.randn(1, 20)
```


```
i2h = torch.mm(W_x, x.t())
h2h = torch.mm(W_h, prev_h.t())
```



# PyTorch Autograd

```
W_h = torch.randn(20, 20, requires_grad=True)
W_x = torch.randn(20, 10, requires_grad=True)
x = torch.randn(1, 10)
prev_h = torch.randn(1, 20)
```

```
i2h = torch.mm(W_x, x.t())
h2h = torch.mm(W_h, prev_h.t())
next_h = i2h + h2h
```






# PyTorch Autograd

```
W_h = torch.randn(20, 20, requires_grad=True)
W_x = torch.randn(20, 10, requires_grad=True)
x = torch.randn(1, 10)
prev_h = torch.randn(1, 20)
```


```
i2h = torch.mm(W_x, x.t())
h2h = torch.mm(W_h, prev_h.t())
next_h = i2h + h2h
```



# PyTorch Autograd

```
W_h = torch.randn(20, 20, requires_grad=True)
W_x = torch.randn(20, 10, requires_grad=True)
x = torch.randn(1, 10)
prev_h = torch.randn(1, 20)
```

```
i2h = torch.mm(W_x, x.t())
h2h = torch.mm(W_h, prev_h.t())
next_h = i2h + h2h
next_h = next_h.tanh()
```

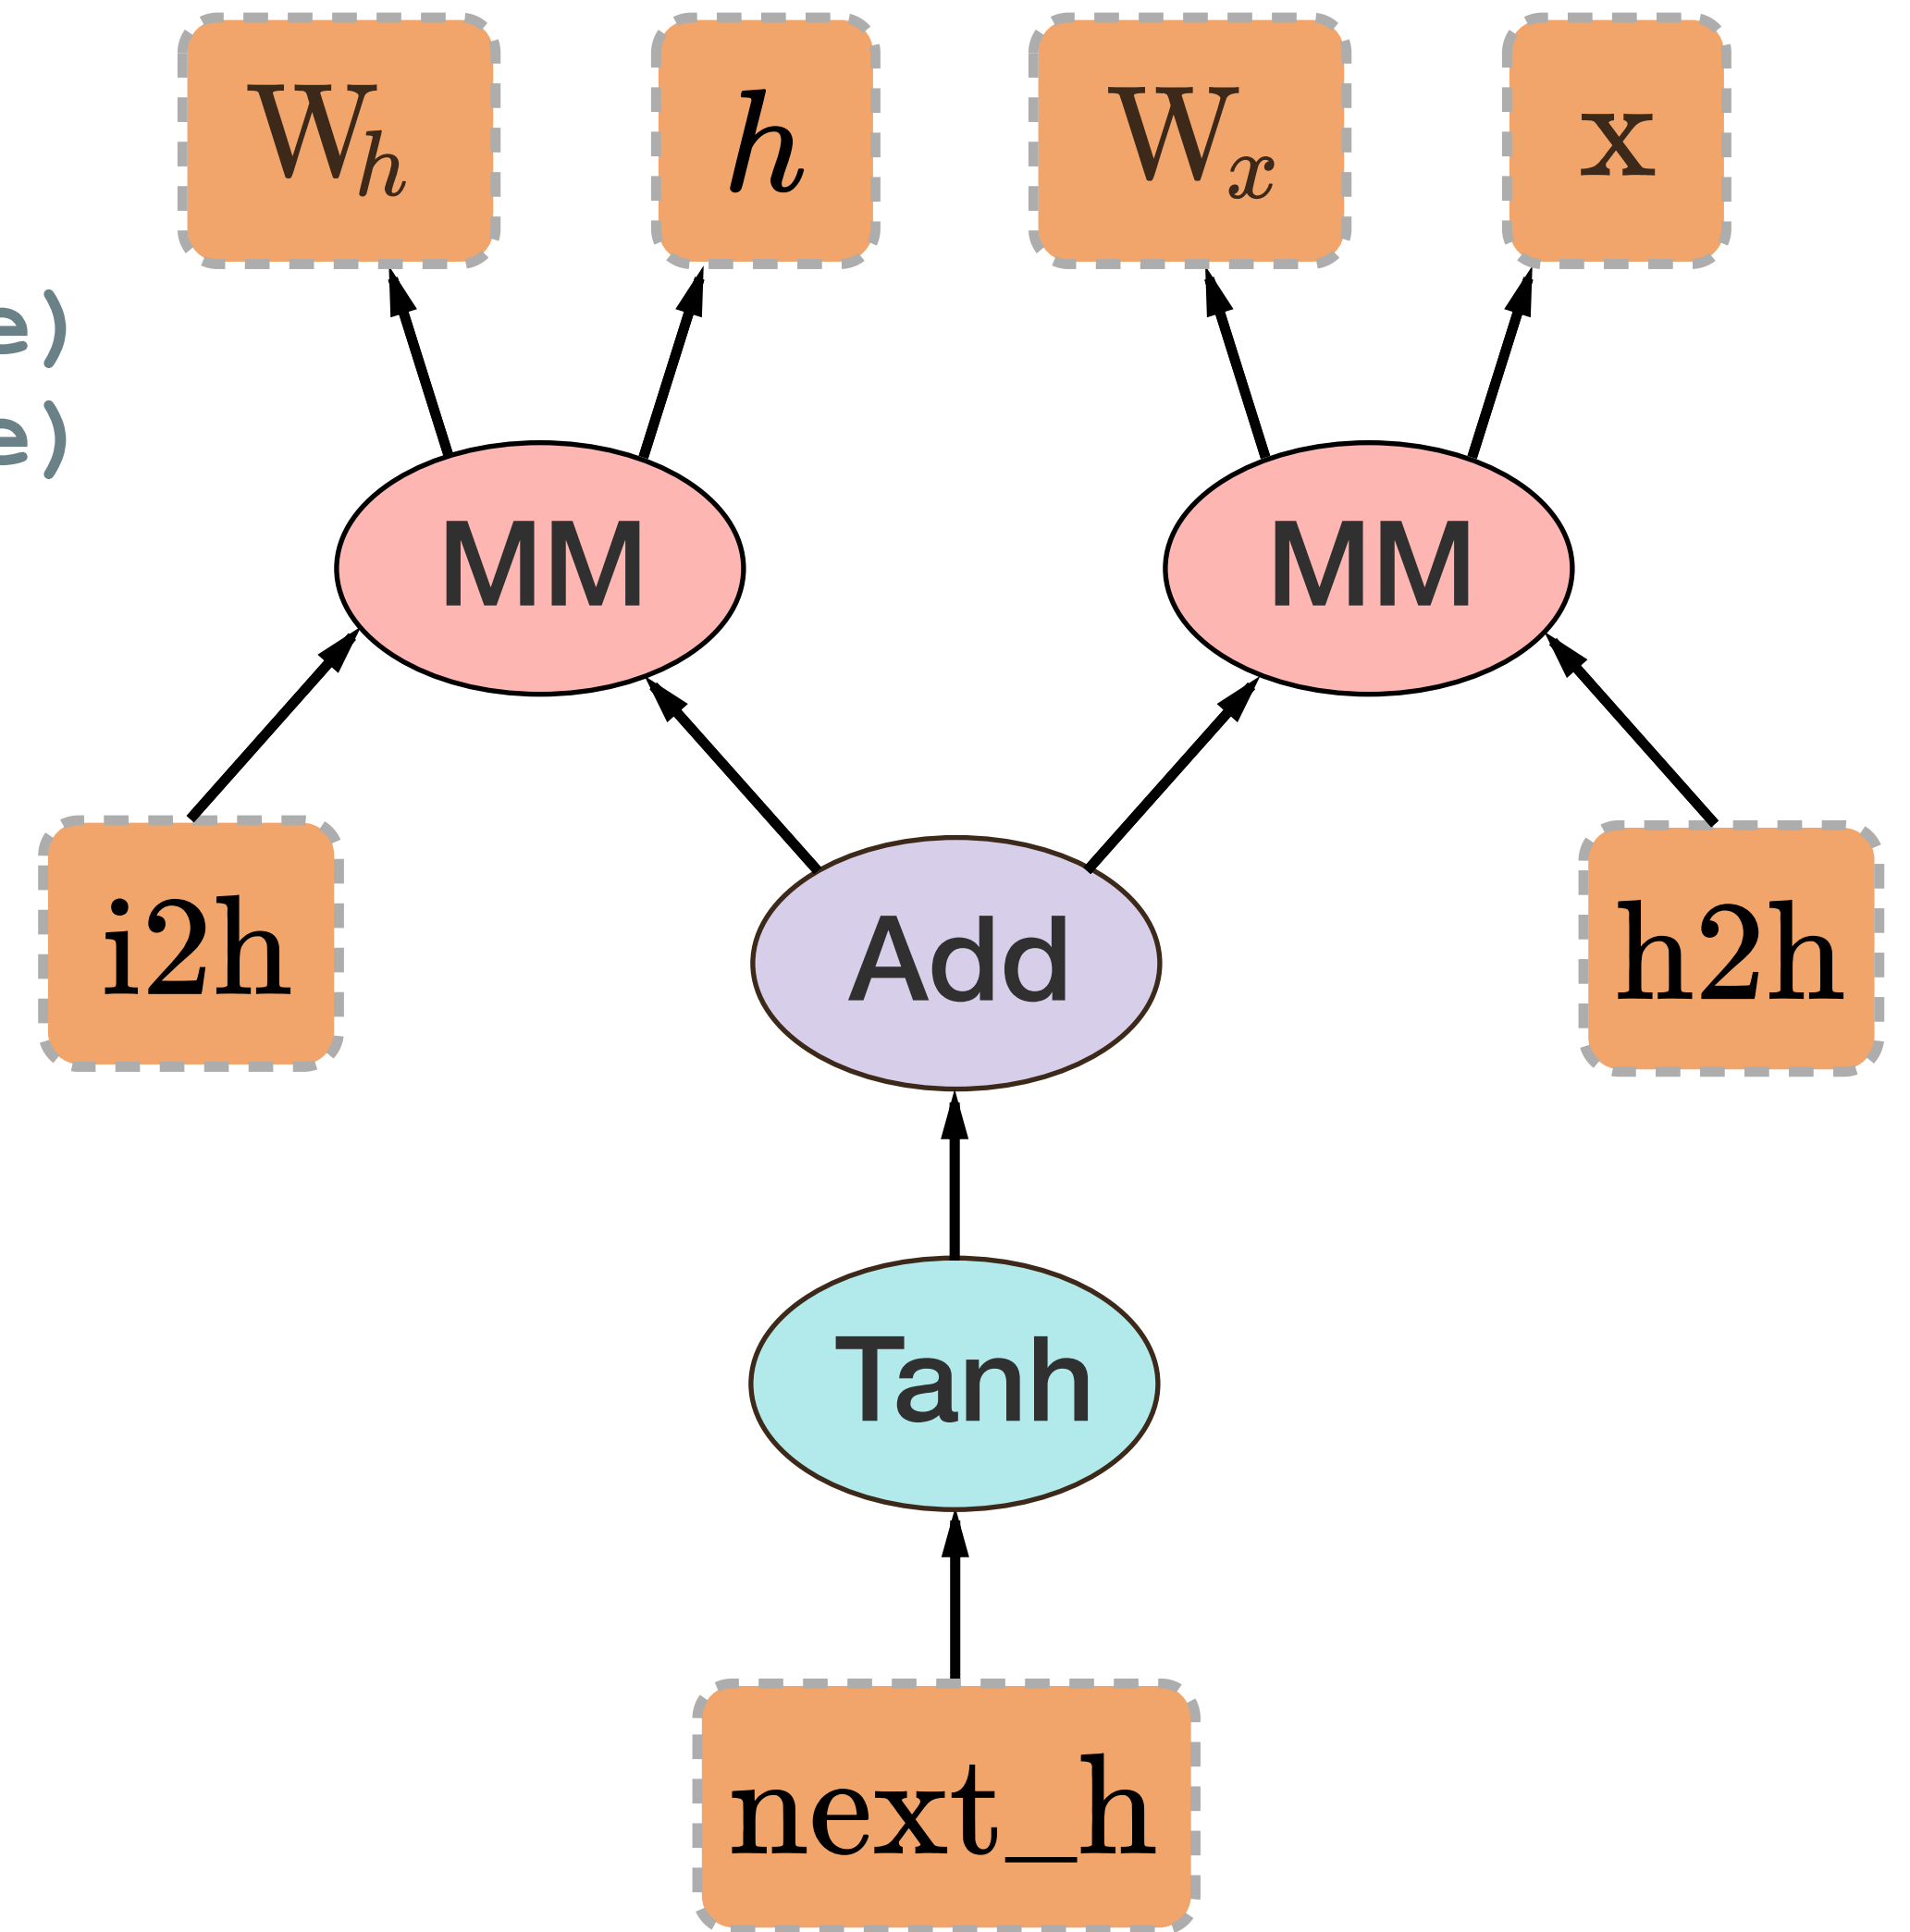


# PyTorch Autograd

```
W_h = torch.randn(20, 20, requires_grad=True)
W_x = torch.randn(20, 10, requires_grad=True)
x = torch.randn(1, 10)
prev_h = torch.randn(1, 20)
```

```
i2h = torch.mm(W_x, x.t())
h2h = torch.mm(W_h, prev_h.t())
next_h = i2h + h2h
next_h = next_h.tanh()
```

```
next_h.backward(torch.ones(1, 20))
```



# Neural Networks

```
1  class Net(nn.Module):
2      def __init__(self):
3          super(Net, self).__init__()
4          self.conv1 = nn.Conv2d(1, 10, kernel_size=5)
5          self.conv2 = nn.Conv2d(10, 20, kernel_size=5)
6          self.conv2_drop = nn.Dropout2d()
7          self.fc1 = nn.Linear(320, 50)
8          self.fc2 = nn.Linear(50, 10)
9
10     def forward(self, x):
11         x = F.relu(F.max_pool2d(self.conv1(x), 2))
12         x = F.relu(F.max_pool2d(self.conv2_drop(self.conv2(x)), 2))
13         x = x.view(-1, 320)
14         x = F.relu(self.fc1(x))
15         x = F.dropout(x, training=self.training)
16         x = self.fc2(x)
17         return F.log_softmax(x)
18
19  model = Net()
20  input = Variable(torch.randn(1, 1, 1, 1))
21  output = model(input)
```

# Neural Networks

```
1  class Net(nn.Module):
2      def __init__(self):
3          super(Net, self).__init__()
4          self.conv1 = nn.Conv2d(1, 10, kernel_size=5)
5          self.conv2 = nn.Conv2d(10, 20, kernel_size=5)
6          self.conv2_drop = nn.Dropout2d()
7          self.fc1 = nn.Linear(320, 50)
8          self.fc2 = nn.Linear(50, 10)
9
10     def forward(self, x):
11         x = F.relu(F.max_pool2d(self.conv1(x), 2))
12         x = F.relu(F.max_pool2d(self.conv2_drop(self.conv2(x)), 2))
13         x = x.view(-1, 320)
14         x = F.relu(self.fc1(x))
15         x = F.dropout(x, training=self.training)
16         x = self.fc2(x)
17         return F.log_softmax(x)
18
19 model = Net()
20 input = Variable(torch.randn(1, 1, 1, 1))
21 output = model(input)
```



# Neural Networks

```
1  class Net(nn.Module):
2      def __init__(self):
3          super(Net, self).__init__()
4          self.conv1 = nn.Conv2d(1, 10, kernel_size=5)
5          self.conv2 = nn.Conv2d(10, 20, kernel_size=5)
6          self.conv2_drop = nn.Dropout2d()
7          self.fc1 = nn.Linear(320, 50)
8          self.fc2 = nn.Linear(50, 10)
9
10     def forward(self, x):
11         x = F.relu(F.max_pool2d(self.conv1(x), 2))
12         x = F.relu(F.max_pool2d(self.conv2_drop(self.conv2(x)), 2))
13         x = x.view(-1, 320)
14         x = F.relu(self.fc1(x))
15         x = F.dropout(x, training=self.training)
16         x = self.fc2(x)
17         return F.log_softmax(x)
18
19 model = Net()
20 input = Variable(torch.randn(1, 1, 1, 1))
21 output = model(input)
```

# Optimization package

SGD, Adagrad, RMSProp, LBFGS, etc.

```
1 net = Net()
2 optimizer = torch.optim.SGD(net.parameters(), lr=0.01, momentum=0.9)
3
4 for input, target in dataset:
5     optimizer.zero_grad()
6     output = model(input)
7     loss = F.cross_entropy(output, target)
8     loss.backward()
9     optimizer.step()
```



# Bootstrapping

Writing  
Dataset loaders

Building models

Implementing  
Training loop

Checkpointing  
models

**Python + PyTorch - an environment to do all of this**

Interfacing with  
environments

Building optimizers

Dealing with  
GPUs

Building  
Baselines



# Bootstrapping

Writing  
Dataset loaders

Building models

Implementing  
Training loop

Checkpointing  
models

## **bootstrapping the Python tooling stack for good UX**

Interfacing with  
environments

Building optimizers

Dealing with  
GPUs

Building  
Baselines



# Python is slow, interpreted

- Global interpreter-lock
- application logic is order of magnitude slower than C++
- moved autograd engine to C++
- moved everything to ATen
  - Side-effect, a clean C++ API

